

Towards a cross-platform, polyglot implementation of *Aggregate Computing* in ScaFi3

Luca Tassinari

25/03/2025

Motivations

- Aggregate Computing span heterogeneous devices and platforms;
- Several implementations of AC exist for different programming languages to:
 - target different platforms and environments;
 - leverage unique strengths of the host programming languages;

However:

- Each of these were developed from scratch, with no code reuse and compatibility in mind;
- No common framework led to fragmentation.

Goal

Investigate the feasibility of building a framework capable of targeting multiple platforms while offering interoperability with other languages.

In particular the work focuses on:

- architectural design of a portable, interoperable layer for Aggregate programming, preserving core abstractions and full code reuse.
- interoperability and distribution strategies enabling seamless data exchange and collective execution across heterogeneous devices and language runtimes;
- evaluation of performance, API idiomacity, and maintenance effort.

⇒ *Scala 3* as the perfect fit to implement AC abstractions and model in a strongly typed internal DSL.

Scala 3 cross-platform capabilities

- **JVM** (desktop, server, Android) & *Java* interop;
- **JS** via Scala.js:
 - Supported platforms:
 - Web (browser);
 - Node.js;
 - WebAssembly (experimental).
 - *JavaScript* interop via annotations, indirectly supporting *TypeScript*;
 - Mature ecosystem.
- **Native** via Scala Native:
 - Supported platforms:
 - `x86-64` and `aarch64` on Linux, macOS and Windows;
 - experimental 32-bit support;
 - suitable for SoC-based IoT devices but not microcontrollers;
 - *C* interop via annotations;
 - Growing ecosystem maturity, limited toolchain support.

ADD PRO AND CONS

Towards a cross-platform polyglot implementation of AC in ScaFi3


Contribution


The contribution of this thesis span three main axes:

1. Add a cross-platform *distribution* module;
2. Add support for a general *cross-platform* and *polyglot* serialization binding;
3. Add a *cross-platform, polyglot* library abstraction layer.

Cross-platform distribution module

- Technology: *stream*, *TCP*-based *connection-oriented sockets*;
 - Each device is bound to a specific *endpoint* (IP + port);
 - Point-to-point connections between neighbors;
 - Neighborhood is statically *fixed* at initialization but can be extended in the future with dynamic discovery strategies;
- Support for multiple platforms: *JVM*, *JS* (Node.js), *Native*;
 - *JVM* + *Native* support via Java Standard *sockets* library;
 - *JS* support via *Node.js net* module using Scala.js type facades;
 - **Implications:**
 - shared code cannot perform blocking operations;
 - all the API is designed to be asynchronous and non-blocking using *Futures*;
 - the primary goal: write as much shared code as possible, minimizing platform-specific implementations.





An example of Scala.js facade over the Node.js `Net` class to allow interoperability with Node.js networking APIs:

```
@js.native
@jsImport("net", JSImport.Namespace)
object Net extends js.Object:

    /** A factory function which creates a new Socket connection. */
    def connect(port: Int, host: String): Socket = js.native

    /** A factory function which creates a new TCP or IPC server. */
    def createServer(connectionListener: js.Function1[Socket, Unit]): Server = js.native
```


Contribution

The contribution of this thesis span three main axes:

1. Add a cross-platform *distribution* module;
2. **Add support for a general *cross-platform* and *polyglot* serialization binding;**
3. Add a *cross-platform*, *polyglot* library abstraction layer.


Cross-platform and polyglot serialization binding

- Devices exchange (ID, Value Tree) pairs
- When exchanging data, values are inserted into the Value Tree encoded using a specific serialization format
 - This is possible since, in the context of an `exchange`, the type information of the value is known
- When receiving data, the Value Tree is decoded but values remain encoded in their serialized format
- Only when the corresponding exchange in the aggregate program is evaluated the value is decoded
 - Again, this is possible since the type information of the expected value is known at that point
- Technically, this is achieved via a combination of Scala 3 *type classes* and *type lambdas* that abstract over the serialization format and allow to cleanly express encoding and decoding requirements as *type bounds*.



```
branch (localId.isEven) {
    nbr(true)
}
{
    nbr(localId)
}
```


Aligned Devices



```
branch (localId.isEven) {
    nbr(true)
}
{
    nbr(localId)
}
```

```
branch (localId.isEven) {
    nbr(true)
}
{
    nbr(localId)
}
```

Aligned
Device 2



- Encodable and Decodable type classes for encoding and decoding generic messages from/to a format (e.g., JSON, binary, ...)

```
/** A type class for encoding messages. */
trait Encodable[-From, +To]:
  /** @return the encoded value in the target type. */
  def encode(value: From): To

  /** A type class for decoding messages. */
  trait Decodable[-From, +To]:
    /** @return the decoded data in the target type. */
    def decode(data: From): To

    /** A type class for encoding and decoding messages. */
    trait Codable[Message, Format] extends Encodable[Message, Format] with Decodable[Format, Message]
    // Type alias for express encodable and decodable capabilities as type bound on values
    type EncodableTo[Format] = [Message] =>> Encodable[Message, Format]
    type DecodableFrom[Format] = [Message] =>> Decodable[Format, Message]
    type CodableFromTo[Format] = [Message] =>> Codable[Message, Format]
```

Every function dealing with, possibly, values distribution add as type bound a Codable instance

```
// inside this function body, Values can be both encoded and decoded
override def xc[Format, Value: CodableFromTo[Format]](
  init: SharedData[Value],
)(
  f: SharedData[Value] => (SharedData[Value], SharedData[Value]),
): SharedData[Value] =
  alignmentScope("exchange"): () =>
    val messages = alignedMessages.map { case (id, value) => id -> value }
    val field = Field(init(localId), messages)
    val (ret, send) = f(field)
    writeValue(send.default, send.alignedValues)
    ret

// extracts the aligned values from the Value Tree and decode them using contextually
// available decoder for Value
def alignedMessages[Format, Value: DecodableFrom[Format]]: Map[DeviceId, Value] = ...

// add a new value into the Value Tree that will be sent to neighbors already serialized using
// contextually available encoder for Value
def writeValue[Format, Value: EncodableTo[Format]](default: Value, overrides: Map[DeviceId, Value]): Unit =
  ...
```

- In non-distribution scenarios, like simulation or local testing, encoding and decoding is a no-op:

```
given forInMemoryCommunications [Message]: Codable[Message, Message] with
  inline def encode(msg: Message): Message = msg
  inline def decode(msg: Message): Message = msg
```

- Useful for API using exchange primitive only for state evolution, like `evolve`, where we do not want to force users to provide encoders/decoders for their values:
 - network managers needs to be implemented to ignore any non-Format values


```
override def evolve[Value](initial: Value)(evolution: Value => Value): Value =
  // `exchange` is called only to update the self-value: `None` is shared with neighbors, so an in-memory
  // codec is enough; non-in-memory network managers will ignore it since it is not serialized.
  exchange(None)(nones =>
    val previousValue = nones(localId).getOrElse(initial)
    nones.set(localId, Some(evolution(previousValue))),
  )(using Codables.forInMemoryCommunications)(localId).get
```

Contribution

The contribution of this thesis span three main axes:

1. Add a cross-platform *distribution* module;
2. Add support for a general *cross-platform* and *polyglot* serialization binding;
3. **Add a *cross-platform, polyglot library abstraction layer*.**

Primary problem: both Scala Native and Scala.js



Towards a cross-platform polyglot implementation of AC in ScaFi3

